

Module 04 Lab - Exploratory Data

Analysis Objective: To learn how to explore a dataset to find patterns, anomalies, and insights before modeling. EDA is like being a detective; you are looking for clues in the data that will help you build a better model. **This lab is fully coded.** Your task is to run each cell, read the detailed explanations, understand the purpose of each visualization, and then complete the experimentation section.

Part 1: Setup and Data Loading**What is Exploratory Data Analysis (EDA)?**EDA is the process of using summary statistics and visualizations to understand a dataset's main characteristics. Before you can build a model, you need to understand your data. What stories does it tell? Are there errors or missing values? Are there strong relationships between variables? EDA helps answer these questions. We will use the famous Titanic dataset for this lab. It contains information about passengers and, crucially, whether they survived the disaster.

```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns # Seaborn is a library built on top of Matplotlib that
```

```
# Load the dataset directly from a URL
df = pd.read_csv('https://raw.githubusercontent.com/datasciencedojo/datasets

print("--- First 5 Rows ---")
print(df.head())

print("--- Basic Info ---")
# .info() is a great first command. It tells us the column names, how many n
# Notice that 'Age' and 'Cabin' have missing values!
df.info()
```

```
--- First 5 Rows ---
  PassengerId  Survived  Pclass  \
0            1         0       3
1            2         1       1
2            3         1       3
3            4         1       1
4            5         0       3

                                Name    Sex  Age  SibSp  \
0                        Braund, Mr. Owen Harris    male  22.0      1
1  Cumings, Mrs. John Bradley (Florence Briggs Th...  female  38.0      1
2                        Heikkinen, Miss. Laina  female  26.0      0
3  Futrelle, Mrs. Jacques Heath (Lily May Peel)  female  35.0      1
4                        Allen, Mr. William Henry    male  35.0      0

   Parch    Ticket   Fare Cabin Embarked
0      0  A/5 21171   7.2500   NaN        S
1      0    PC 17599  71.2833   C85        C
2      0  STON/O2. 3101282   7.9250   NaN        S
3      0    113803  53.1000  C123        S
4      0   373450   8.0500   NaN        S

--- Basic Info ---
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
#   Column          Non-Null Count  Dtype
---  -
0   PassengerId     891 non-null    int64
1   Survived        891 non-null    int64
2   Pclass          891 non-null    int64
3   Name            891 non-null    object
4   Sex             891 non-null    object
5   Age             714 non-null    float64
6   SibSp           891 non-null    int64
7   Parch           891 non-null    int64
8   Ticket          891 non-null    object
9   Fare            891 non-null    float64
10  Cabin           204 non-null    object
11  Embarked        889 non-null    object
dtypes: float64(2), int64(5), object(5)
memory usage: 83.7+ KB
```

Part 2: Descriptive Statistics Let's start by getting a high-level numerical summary of the data. The `.describe()`

- method is perfect for this. It calculates statistics like mean, standard deviation, min, and max for the numerical columns.

```
# Get summary statistics for numerical columns
print("--- Descriptive Statistics ---")
print(df.describe())
print("--- Key Insights from Statistics ---")
print(f"The average age of a passenger was {df['Age'].mean():.1f} years.")
print(f"The overall survival rate was {df['Survived'].mean():.1%}.")
print(f"Fares ranged from ${df['Fare'].min()} to a whopping ${df['Fare'].max()}")
```

```
--- Descriptive Statistics ---
```

	PassengerId	Survived	Pclass	Age	SibSp	\
count	891.000000	891.000000	891.000000	714.000000	891.000000	
mean	446.000000	0.383838	2.308642	29.699118	0.523008	
std	257.353842	0.486592	0.836071	14.526497	1.102743	
min	1.000000	0.000000	1.000000	0.420000	0.000000	
25%	223.500000	0.000000	2.000000	20.125000	0.000000	
50%	446.000000	0.000000	3.000000	28.000000	0.000000	
75%	668.500000	1.000000	3.000000	38.000000	1.000000	
max	891.000000	1.000000	3.000000	80.000000	8.000000	

	Parch	Fare
count	891.000000	891.000000
mean	0.381594	32.204208
std	0.806057	49.693429
min	0.000000	0.000000
25%	0.000000	7.910400
50%	0.000000	14.454200
75%	0.000000	31.000000
max	6.000000	512.329200

```
--- Key Insights from Statistics ---
The average age of a passenger was 29.7 years.
The overall survival rate was 38.4%.
Fares ranged from $0.0 to a whopping $512.3292.
```

Part 3: Visual EDA - Telling Stories with Plots Numbers are great, but plots make patterns and relationships immediately obvious. The goal of visual EDA is to turn data into insights. **A Note on Plotting Libraries:** * **Matplotlib:**

- ✓ The foundational library, gives you full control over everything.* **Seaborn**: Built on Matplotlib, it provides a simpler, high-level interface for creating common statistical plots. We will use Seaborn for its ease of use and attractive defaults.

Visualization 1: How many survived?A simple `countplot` is the best


- ✓ way to see the distribution of a categorical variable, like our target `Survived`.


```
sns.set_style('whitegrid') # Sets a nice visual style for our plotsplt.figure
```

Visualization 2: Does passenger class matter for survival?Now we

- ✓ want to see if there's a relationship between two variables. We can use the `hue` parameter in `countplot` to split the bars by another category.

```
plt.figure(figsize=(10, 6))
sns.countplot(x='Pclass', hue='Survived', data=df)
plt.title('Survival by Passenger Class')
plt.legend(['Died', 'Survived'])
plt.show()
print("Insight: This is a very strong pattern. 1st class passengers had a mu
```






Insight: This is a very strong pattern. 1st class passengers had a much higher

- ✓ Visualization 3: What about gender? The 'women and children first' mantra is famous. Let's see if the data supports it.


```
plt.figure(figsize=(10, 6))
sns.countplot(x='Sex', hue='Survived', data=df)
plt.title('Survival by Gender')
plt.legend(['Died', 'Survived'])
plt.show()
print("Insight: The pattern is undeniable. A much higher proportion of females survived.")
```



Insight: The pattern is undeniable. A much higher proportion of females survived.

Visualization 4: How does age play a role? For a continuous variable like `Age`, a histogram is a great way to see its distribution.

```
# A FacetGrid allows us to create multiple plots side-by-side to compare dis  
# Here, we create one histogram for passengers who died (col='Survived'=0) a  
g = sns.FacetGrid(df, col='Survived', height=6)  
g.map(plt.hist, 'Age', bins=20)  
plt.show()  
print("Insight: The age distribution for those who did not survive is center
```




Insight: The age distribution for those who did not survive is centered around

Part 4: Student Experimentation **Instructions:** Create your own visualizations to explore other relationships in the data.

Experiment 1: Port of Embarkation1. The `Embarked` column tells you where a passenger boarded the ship (C = Cherbourg, Q =


Queenstown, S = Southampton).2. Create a `countplot` to see how survival rates differed by the port of embarkation. Does where they boarded seem to be related to their survival?

```
# --- ENTER YOUR CODE HERE ---  
countplot = sns.countplot(x='Embarked', hue='Survived', data=df)  
plt.title('Survival by Port of Embarkation')  
plt.legend(['Died', 'Survived'])  
plt.show()
```



Experiment 2: Fare vs. Survival1. `Fare` is a continuous numerical variable.2. A `boxplot` or `violinplot` is a great way to see the distribution of fares for those who survived vs. those who didn't.3. Create one of these plots to compare the `Fare` distribution by `Survived`. What does it tell you?

```
# --- ENTER YOUR CODE HERE ---  
sns.boxplot(x='Survived', y='Fare', data=df)  
plt.title('Fare Distribution by Survival')  
plt.show()
```



Knowledge Check Instructions: Answer the following questions in this markdown cell.

1. **What is the primary goal of Exploratory Data Analysis (EDA)** -- EDAs allow for visualization of data to help quickly

determine correlatives and outliers for data scrubbing.

2. Based on the plots in this lab, what kind of person had the best chance of surviving the Titanic? (Describe them in terms of class, gender, and age). -- High class persons, females, and children had the best chances of survival according to the data.

3. Why is it important to visualize data instead of just looking at summary statistics? What can a plot show you that a number like 'mean' or 'count' can't? -- It can show you categorial trends.